AI ENGINEERING | FULL-TIME

AI Engineer

Location: Bengaluru / Hybrid | Experience: 0-2 years | Employment: Full-time

Role overview. We are looking for an AI Engineer to build and deploy production-ready AI applications using large language models, retrieval-augmented generation (RAG), APIs, vector databases, and Python backends. You will work across experimentation, evaluation, backend engineering, and cloud deployment to turn AI capabilities into reliable user-facing products.

What you will do

- Build AI-powered applications using Python, LLM APIs, prompt engineering, structured outputs, and tool/function calling.
- Design and implement RAG pipelines including document ingestion, chunking, embeddings, vector search, reranking, context assembly, and response generation.
- Develop backend services and APIs using frameworks such as FastAPI or Flask, with authentication, validation, logging, error handling, and rate limiting.
- Integrate vector databases and data stores such as Qdrant, FAISS, PostgreSQL, or Redis as appropriate for the application.
- Evaluate AI systems using retrieval and generation metrics, test sets, latency/cost measurements, and qualitative error analysis.
- Create agents or workflow-based systems that can use tools, follow multi-step processes, and maintain conversational context.
- Containerize applications with Docker and support deployment to cloud platforms such as AWS, GCP, or Azure.
- Collaborate with software engineers and product teams to translate requirements into maintainable, observable AI services.

Required skills

Area	Expected skills
Programming	Strong Python fundamentals; OOP; modules/packages; virtual environments; Git; debugging.
LLMs	LLM APIs; prompt design; system/user messages; structured JSON outputs; function/tool calling; token and context management.
RAG	Embeddings; vector similarity; chunking; metadata filtering; retrieval; context construction; basic RAG evaluation.
Backend	REST APIs; FastAPI or Flask; Pydantic; authentication basics; request validation; exception handling.
Data	SQL fundamentals; PostgreSQL or equivalent; basic pandas/NumPy familiarity.

Area	Expected skills
Deployment	Docker; environment variables and secrets; basic CI/CD; familiarity with a cloud platform and logs/monitoring.

Preferred skills

- LangChain, LlamaIndex, LangGraph, MCP, or equivalent AI application frameworks.
- Qdrant, Pinecone, Weaviate, Milvus, Chroma, or advanced FAISS usage.
- Async Python, background jobs, queues, caching, and performance optimization.
- LLM observability/evaluation tools, prompt/version management, guardrails, and AI security practices.
- Hands-on experience with OpenAI, Google Gemini, Anthropic, Groq, or open-source models.
- Basic ML concepts including classification, regression, embeddings, model evaluation, and data preprocessing.

Qualifications

- Bachelor's degree in Computer Science, Engineering, Data Science, or a related technical field, or equivalent practical experience.
- Demonstrated ability to build at least one end-to-end AI project with a working API and user-facing interface.
- Comfortable reading documentation, debugging dependency/version issues, and learning unfamiliar APIs quickly.

Example project scope

Typical first project: Build an AI assistant that accepts PDFs or web content, indexes the content in a vector database, retrieves relevant context, answers user questions with citations, stores conversation history, exposes a FastAPI backend, and is deployed in Docker on a cloud platform. Add authentication, rate limiting, logging, automated tests, and an evaluation dataset.

What we value

- Production mindset: reliability, security, cost awareness, maintainability, and clear failure handling.
- Curiosity and ownership: ability to investigate technical problems independently and communicate trade-offs clearly.
- Evidence of learning: GitHub projects, technical write-ups, internships, coursework, or open-source contributions.

Interview focus: Python, APIs/backend fundamentals, LLM concepts, RAG architecture, vector search, debugging, deployment basics, and discussion of personal AI projects.